

CS444/544

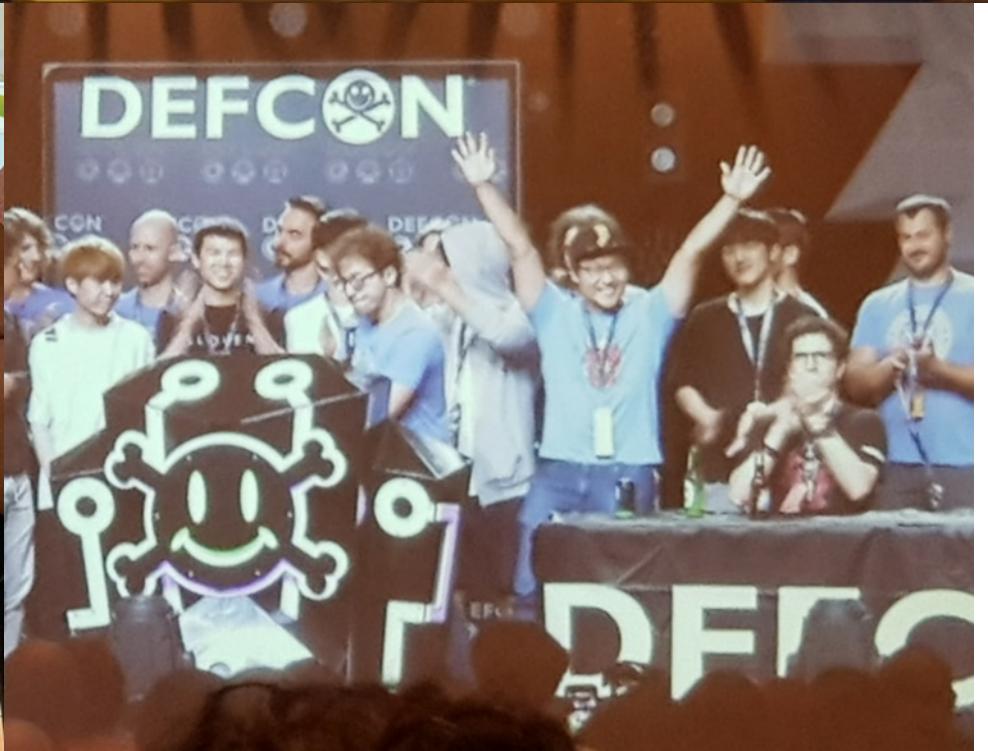
Operating Systems II

Course Intro

09/24/20



Oregon State
University



Instructor: Yeongjin Jang

- At OSU Since Oct 2017
- Main research area: systems security
- Hacking
 - CPU side-channel attacks, Jailbreaking, exploit development, automatic hacking (fuzzing / symbolic execution), designing secure systems, etc..
- Faculty advisor of OSUSEC
- CS419/579 Cyber Attacks and Defense
- CS419/579 Systems Security
- Feel free to reach me if you are interested in cybersecurity...
 - Join OSUSEC!

```

static void *
boot_alloc(uint32_t n)
{
    static char *nextfree; // virtual
    char *result;

    // Initialize nextfree if this is
    // 'end' is a magic symbol automatically
    // which points to the end of the
    // the first virtual address that
    // to any kernel code or global
    if (!nextfree) {
        extern char end[];
        nextfree = ROUNDUP((char *)end);
    }
}

```

```

static inline physaddr_t
page2pa(struct PageInfo
{
    return (pp - pages)
}

// These variables are set in
pde_t *kern_pgd; // Kernel
struct PageInfo *pages; // Page table
static struct PageInfo *page_f
Device drivers
Applications
}

```

```

#define TRAPHANDLER_NOEC(name, num)
.globl name;
.type name, @function;
.align 2;
name:
pushl $0;
pushl $(num);
jmp _alltraps

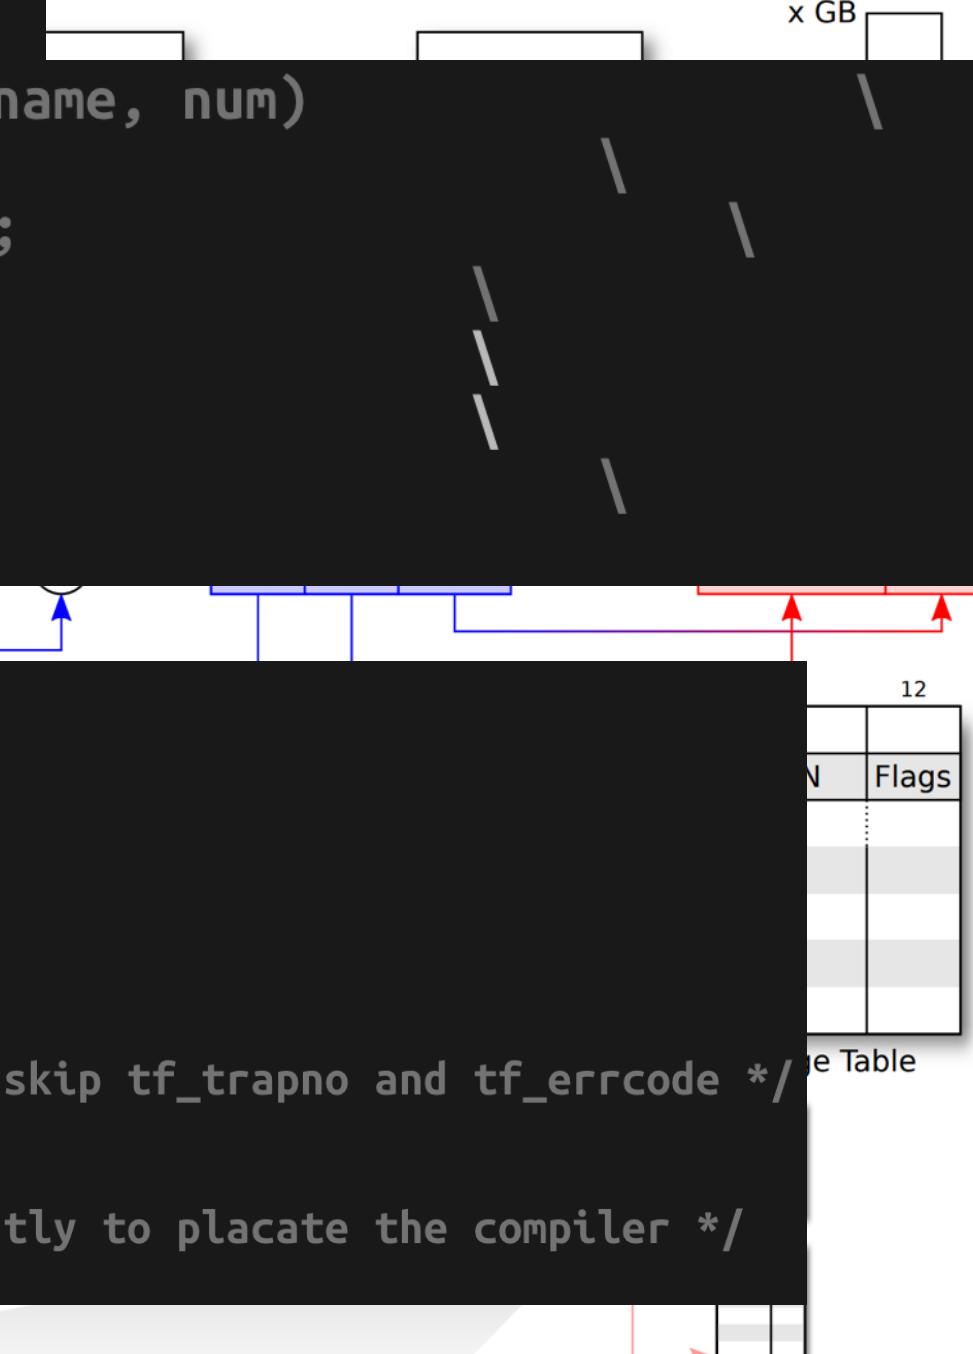
```

```

void
env_pop_tf(struct Trapframe *tf)
{
    asm volatile(
        "\tmovl %0,%%esp\n"
        "\tpopal\n"
        "\tpopl %%es\n"
        "\tpopl %%ds\n"
        "\ttaddl $0x8,%%esp\n" /* skip tf_trapno and tf_errcode */
        "\tiret\n"
        : : "g" (tf) : "memory");
    panic("iret failed"); /* mostly to placate the compiler */
}

```

d-Mode Address Translation



Course Objective

- Understand how modern computer systems work (in detail)
- Be able to answer the following questions:
 - What happens when we **turn on** the computer? How does it **boot**?
 - How an OS **runs an application**?
 - How an OS runs application that requires **memory more than its physical memory**?
 - How **multiple applications can run** on the system?
 - How an OS enforces **privilege separation**?
 - How an OS protects itself from **malicious software**?
 - How multiple programs **synchronize** each other? How can we **implement a lock**?

Important Links

- Website: <https://os.unexploitable.systems>
- Instructor: Dr. Yeongjin Jang (yeongjin.jang@oregonstate.edu)
- TAs:
 - Hadi Rahal-Arabi, Andrew Quach, and Phillip Mestas III (graduate TAs)
 - Laura Jiang, Lyell Read, and Corbin Tegner (undergraduate TAs)
- Gitlab: <https://gitlab.unexploitable.systems>
- Piazza: <https://piazza.com/class/ket6bdq5r6y55>
- Discord: <https://discord.gg/hPyPaQf>
- Assignment server: os2.engr.oregonstate.edu, os1, oldos1, oldos2

Course Structure

- 10 weeks schedule
 - <https://os.unexploitable.systems/cal.html>
 - Virtualization (Week 1-4)
 - Concurrency (Week 5-8)
 - Persistency and others (Week 9-10)
- Textbook
 - <http://pages.cs.wisc.edu/~remzi/OSTEP/>
- **Read:** prep materials posted on homepage
- **Watch:** 50 min VIDEO lecture
- **Study:** study JOS labs (tutorial videos / lab instructions website)
- **Engage:** ~30 min ZOOM Q&A and office hours on Discord, discuss with peers!

Monday	Tuesday	Wednesday	Thursday	Friday
Sep 21	Sep 22	Sep 23	Sep 24 LEC 1: Intro to the course Watch 1: Lecture #1 VIDEO PDF PPTX Watch 2: Lab tutorial 1 VIDEO PDF PPTX Study, Lab 1: Booting a PC Read: Textbook Read: at&t_asm GDB tutorial1 tutorial2 cheat-sheet Read: tmux cheatsheet (ctrl-b -> backtick) tmux-cheat-sheet First day of class	Sep 25
Sep 28	Sep 29 LEC 2: BIOS/Booting/CPU	Sep 30	Oct 1 LEC 3: Memory: Address Space, Segmentation, and Paging Read: x86_Address_Translation Read: Textbook1 Textbook2 Textbook3	Oct 2
Oct 5 DUE: Lab 1 (100%)	Oct 6 LEC 4: Virtual Address Translation Study, Lab 2: Memory Management Read: Page_Table	Oct 7	Oct 8 LEC 5: Virtual Memory Layout Read: Textbook1 Textbook2 Textbook3 Textbook4 Textbook5 Textbook6	Oct 9
Oct 12 DUE: Lab 1 (75%)	Oct 13 LEC 6: JOS Memory Management	Oct 14	Oct 15 Quiz 1: Virtual Memory	Oct 16

Meeting Time (with me)

- Lecture (Video, **asynchronous**)
 - Released at Tu/Th 12:00pm
 - Video, link will be available on Canvas/Homepage
- Lecture Q&A (**Synchronous** via Zoom, **not required**)
 - Tu/Th 01:00pm ~ 01:30pm (Zoom link on Canvas)
- My office hour (via Discord)
 - Wed 06:00pm – 07:30pm (<https://discord.gg/hPyPaQf>)

Recitation / TA Office Hours

- Available via Discord
 - <https://discord.gg/hPyPaQf>
- We will use recitation section times as TA office hour as well
 - Total 28.5 hours of office hours
 - Monday: 4 hours
 - Tuesday: 6 hours
 - Wednesday: 7.5 hours
 - Thursday: 5 hours
 - Friday: 6 hours

A screenshot of a Discord message thread. The first message is from a user named 'Hi Professor' on April 22, 2020, asking about the `page_alloc()` function. The user has provided some code and asks where to use `page2kva`. A TA named 'yeongjin' responds, explaining that `free_page` is an object of `struct PageInfo` and that `page2kva` is used to initialize the physical page content to 0, not the `PageInfo` object. The TA provides the code for `memset(page2kva(free_page), 0, PGSIZE)`.

04/22/2020
Hi Professor

I have a question about page_alloc()
This is what I have

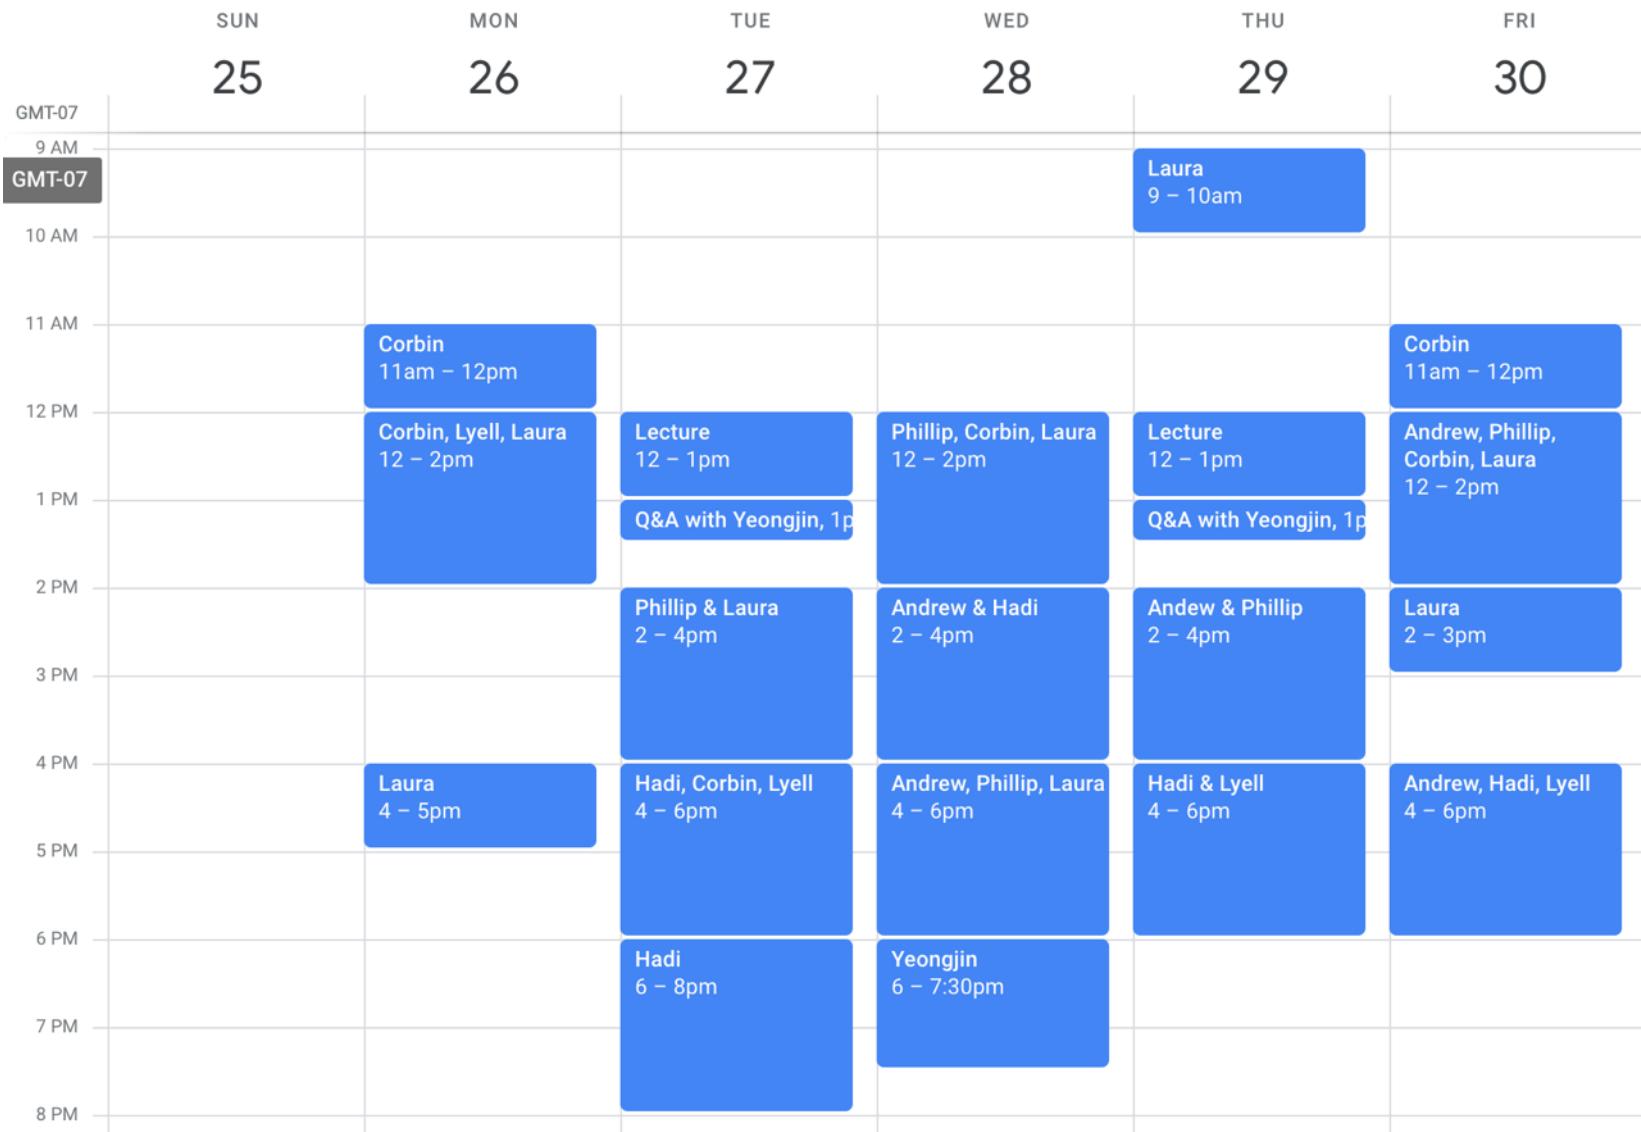
```
301 // Hint: use page2kva and memset
302 struct PageInfo *
303 page_alloc(int alloc_flags)
304 {
    // Fill this function in
305
    // If out of free memory, return NULL
306    if (page_free_list == NULL)
307    {
        return NULL;
    }
    // Cut the head, return the addr, reconnect head
308    struct PageInfo *free_page = page_free_list;
309    page_free_list = free_page->pplink;
310    free_page->pplink = NULL;
311
312    if (alloc_flags & ALLOC_ZERO == 1)
313    {
        memset(free_page, 0, PGSIZE);
    }
314
315    return free_page;
316
317
318
319
320
321
322
323
324
325
```

I'm not sure where to use page2kva
it looks like i need to return a pointer to struct PageInfo, which I did
so what's page2kva for?

yeongjin 04/22/2020

free_page is an object of struct PageInfo.
And you might want to initialize the content of the corresponding physical page to 0, not the struct PageInfo object.
so basically what you should do is
`memset(page2kva(free_page), 0, PGSIZE)`
`free_page -> physical address of the page represented by free_page -> virtual address of that physical address -> then you can use memset`

Office Hours, etc.



<https://calendar.google.com/calendar/u/0/embed?src=06a8uv6u8f01gkg4nljsc35kqg@group.calendar.google.com&ctz=America/LosAngeles&mode=WEEK>

Move to the week of 9/25~10/2

Grading

- 70% JOS lab assignment
 - Lab 1 (10%), Lab 2 (15%), Lab 3 (20%), Lab 4 (25%)
- 30% Quizzes (mini-exam)
 - Quiz 1 (10/15) : Virtual Memory
 - Quiz 2 (11/05): System calls, faults, and exceptions
 - Quiz 3 (11/19): Concurrency

Grading Scheme (tentative):

100 >= A >= 93 (96 for graduate students)
93 > A- >= 90 (93)
90 > B+ >= 86 (89)
86 > B >= 83 (86)
83 > B- >= 80 (83)
80 > C+ >= 76 (79)
76 > C >= 73 (76)
73 > C- >= 70 (73)
70 > D+ >= 66 (69)
66 > D >= 63 (66)
63 > D >= 60 (63)
F < 60 (63)

The Lab (70%)

- Four labs
 - JOS Lab 1 (10%): Booting a PC (1.5 weeks, due on 10/5)
 - Bootloader, protected mode, etc.
 - JOS Lab 2 (15%): Memory Management (2 weeks, due on 10/19)
 - Virtual memory, paging, etc.
 - JOS Lab 3 (20%): User Environment (3 weeks, due on 11/9)
 - Process, user, kernel, system call, etc.
 - JOS Lab 4 (25%): Preemptive Multitasking (3.5 weeks, due on 12/4)
 - Implementing context switching, multi-core support, inter-process communication, etc.

How to Conduct Lab Assignments?

- Visit Lab Tutorial Webpage
 - <https://os.unexploitable.systems/lab/lab1.html>
- Watch Lab Tutorial Video
 - I will explain necessary stuff for the lab assignments in the video (code/examples, etc.) and also share some tips...

An Exercise Example in Lab 1

Note

Exercise 3. Take a look at the [lab tools guide](#), especially the section on GDB commands. Even if you're familiar with GDB, this includes some esoteric GDB commands that are useful for OS work.

Set a breakpoint at address 0x7c00, which is where the boot sector will be loaded. Continue execution until that breakpoint. Trace through the code in `boot/boot.S`, using the source code and the disassembly file

`obj/boot/boot.asm` to keep track of where you are. Also use the `x/i` command in GDB to disassemble sequences of instructions in the boot loader, and compare the original boot loader source code with both the disassembly in `obj/boot/boot.asm` and GDB.

Trace into `bootmain()` in `boot/main.c`, and then into `readsect()`. Identify the exact assembly instructions that correspond to each of the statements in `readsect()`. Trace through the rest of `readsect()` and back out into `bootmain()`, and identify the begin and end of the `for` loop that reads the remaining sectors of the kernel from the disk. Find out what code will run when the loop is finished, set a breakpoint there, and continue to that breakpoint. Then step through the remainder of the boot loader.

The Lab Could be



- Coding KERNEL code in C
 - Any memory error -> Triple fault

Intel® 64 and IA-32 Architectures
Manual

- Us

Watch both lectures and lab tutorial
videos on time and ask TAs for help!

- As

- Co

- Page table
- Global descriptor table (GDT)
- Interrupt descriptor table (IDT)

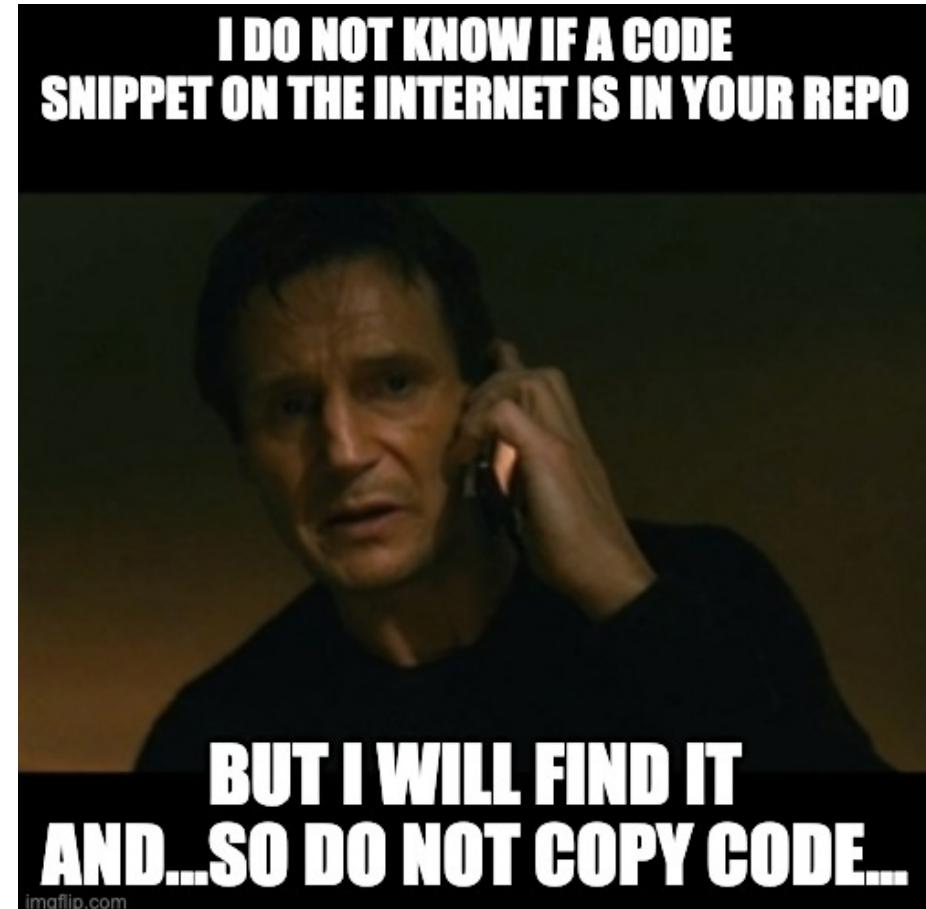
		IA-32e	CR3	47:55	N/A (PS must be 0)
Page-directory-pointer table	PDPTE	32-bit			N/A
		PAE	CR3	31:30	N/A (PS must be 0)
		IA-32e	PML4E	38:30	1-GByte page if PS=1 ¹
Page directory	PDE	32-bit	CR3	31:22	4-MByte page if PS=1 ²
		PAE, IA-32e	PDPTE	29:21	2-MByte page if PS=1
Page table	PTE	32-bit		21:12	4-KByte page
		PAE, IA-32e	PDE	20:12	4-KByte page

Lab Rules

- DO NOT SHARE YOUR CODE WITH OTHER STUDENTS
 - You are encouraged to discuss with others about the assignments but do not ask/give the code to the others
 - Do not copy other students' code or code available in online
 - Do not publish your code online
- You will be asked to submit a simple write-up for the assignment
 - Describe how you solve each exercise/questions
 - Mention your collaborators in the write-up
 - Do not copy other students' write-up
 - Do not publish your write-up online

Lab Rules

- Plagiarism will be punished via the Office of Student Life..
 - E.g., getting F or zero point for the lab assignment that matters with plagiarism...
- Please refer the Code of Student Conduct
 - <https://studentlife.oregonstate.edu/studentconduct/academicmisconduct>
 - [https://studentlife.oregonstate.edu/sites/studnetlife.oregonstate.edu/files/edited code of student conduct.pdf](https://studentlife.oregonstate.edu/sites/studnetlife.oregonstate.edu/files/edited_code_of_student_conduct.pdf)



Lab Rules – Late Submissions

- If you submit your assignment before the due date, then
 - You will get 100% of credit from ‘make grade’ of your submission
- If you submit your assignment within one week after the due date, then
 - You will get 75% of credit from ‘make grade’ of your submission
- If you submit your assignment before 12/11 11:59 pm, then
 - You will get 50% of credit from ‘make grade’

Due Dates on the Calendar

Oct 5 DUE: Lab 1 (100%)	Oct 6 LEC 4: Virtual Address Translation Study, Lab 2: Memory Management Read: Page_Table	Oct 7	Oct 8 LEC 5: Virtual Memory Layout Read: Textbook1 Textbook2 Textbook3 Textbook4 Textbook5 Textbook6	Oct 9
Oct 12 DUE: Lab 1 (75%)	Oct 13 LEC 6: JOS Memory Management	Oct 14	Oct 15 Quiz 1: Virtual Memory	Oct 16
Oct 19 DUE: Lab 2 (100%)	Oct 20 LEC 7: Quiz 1 Review Study, Lab 3: User Environment	Oct 21	Oct 22 LEC 8: User/Kernel Switch Read: Textbook-process Textbook-syscall Textbook-trap	Oct 23
Oct 26 DUE: Lab 2 (75%)	Oct 27 LEC 9: Handling Interrupts/Exceptions	Oct 28	Oct 29 LEC 10: System Calls and Page Fault	Oct 30
Nov 2	Nov 3 LEC 11: Virtualization Recap and Quiz 2 Prep	Nov 4	Nov 5 Quiz 2: System calls, faults, and exceptions	Nov 6
Nov 9 DUE: Lab 3 (100%)	Nov 10 LEC 12: Multi-threading and Synchronization Study, Lab 4: Preemptive Multitasking Read: READ Concurrency Thread	Nov 11	Nov 12 LEC 13: Lock and Thread Synchronization	Nov 13
Nov 16 DUE: Lab 3 (75%)	Nov 17 LEC 14: Concurrency Bugs and Deadlock Read: READ Bugs	Nov 18	Nov 19 Quiz 3: Concurrency	Nov 20
Nov 23 Thanksgiving break	Nov 24 Thanksgiving break	Nov 25 Thanksgiving break	Nov 26 Thanksgiving break	Nov 27 Thanksgiving break
Nov 30	Dec 1 LEC 15: TBD	Dec 2	Dec 3 LEC 16: Final Summary <i>The last day of class</i>	Dec 4 DUE: Lab 4 (100%)
Dec 7 No Final Exam!	Dec 8 No Final Exam!	Dec 9 No Final Exam!	Dec 10 No Final Exam!	Dec 11 DUE: Lab 1,2,3 (50%), 4 (75%)

CS 544 Students

- Will have higher grade bar than CS444 (+3pts)
 - E.g., A is for 93 and over for CS444, and 96 and over for CS544

Grading Scheme (tentative):

100 >= A >= 93 (96 for graduate students)

93 > A- >= 90 (93)

90 > B+ >= 86 (89)

86 > B >= 83 (86)

83 > B- >= 80 (83)

80 > C+ >= 76 (79)

76 > C >= 73 (76)

73 > C- >= 70 (73)

70 > D+ >= 66 (69)

66 > D >= 63 (66)

63 > D >= 60 (63)

F < 60 (63)

Tips to the Lab

- Study in a group (**discussions are highly encouraged!**)
 - But please **write the code individually!**
- Follow tutorial video
- Ask questions (Piazza, Discord)
- Understand your time budget (debugging will take **lots of your time!**)
 - **Plan ahead** to finish the labs on time
- Learn basic tools (e.g., C, gdb, assembly, editors, tmux, etc.) ASAP
 - This will help you earn more time on doing labs...
 - <https://missing.csail.mit.edu/>
 - Up to Debugging and Profiling would be helpful...

Notices

- There will be lecture Q&A from 1:00 pm to 1:30 pm on every Tuesday and Thursday
- My (virtual) office hour is Wed 6:00 pm ~ 7:30 pm (Discord)